

LSTM and Named Entity Recognition


- RNNs and Vanishing Gradients
- LSTMs
- Named Entity Recognition

RNNs and Vanishing Gradients

- RNNs
 - Backprop through time
 - RNNs and vanishing/exploding gradients
 - Solutions
- Advantages
 - Captures dependencies within a short range
 - Takes up less RAM than other n-gram models
- Disadvantages
 - Struggles to capture long term dependencies
 - Prone to vanishing or exploding gradients

RNN Basic Structure

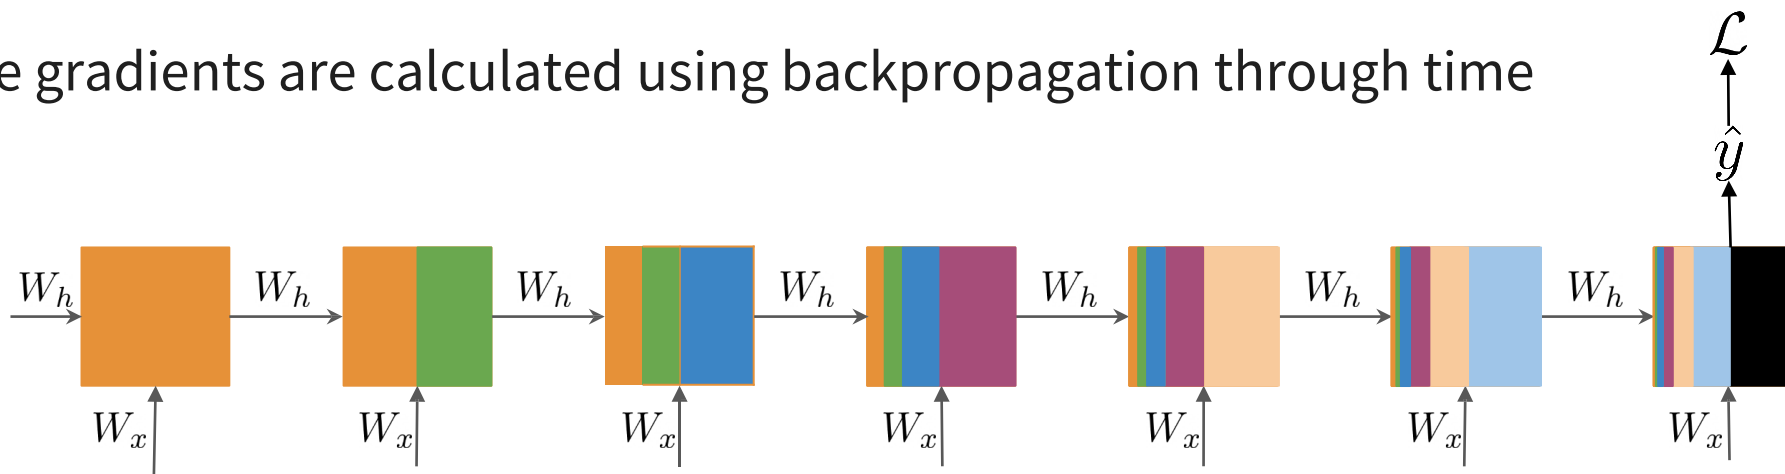
She kicked the ball right into the _____



Learnable parameters

Backpropagation through time

- The gradients are calculated using backpropagation through time



W_x
 W_h

Same at every step

$$\frac{\partial L}{\partial W_h} \propto \sum_{1 \leq k \leq t} \left(\prod_{t \geq i > k} \frac{\partial h_i}{\partial h_{i-1}} \right) \frac{\partial h_k}{\partial W_h}$$

Gradient is proportional to a sum of partial derivative products

Backpropagation through time

$$\frac{\partial L}{\partial W_h} \propto \sum_{1 \leq k \leq t} \left(\prod_{t \geq i > k} \frac{\partial h_i}{\partial h_{i-1}} \right) \frac{\partial h_k}{\partial W_h} \rightarrow \text{Contribution of hidden state } k$$

Length of the product proportional to
how far k is from t

$$\frac{\partial h_t}{\partial h_{t-1}} \frac{\partial h_{t-1}}{\partial h_{t-2}} \frac{\partial h_{t-2}}{\partial h_{t-3}} \frac{\partial h_{t-3}}{\partial h_{t-4}} \frac{\partial h_{t-4}}{\partial h_{t-5}} \frac{\partial h_{t-5}}{\partial h_{t-6}} \frac{\partial h_{t-6}}{\partial h_{t-7}} \frac{\partial h_{t-7}}{\partial h_{t-8}} \frac{\partial h_{t-8}}{\partial h_{t-9}} \frac{\partial h_{t-9}}{\partial h_{t-10}} \frac{\partial h_{t-10}}{\partial W_h}$$

Contribution of hidden state $t-10$

- The contribution to the gradient of a hidden state that is 10 steps away from step t

Backpropagation through time

$$\frac{\partial L}{\partial W_h} \propto \sum_{1 \leq k \leq t} \left(\prod_{t \geq i > k} \frac{\partial h_i}{\partial h_{i-1}} \right) \frac{\partial h_k}{\partial W_h}$$

Contribution of hidden state k

Length of the product proportional to
how far k is from t

Partial derivatives < 1

Contribution goes to 0

Vanishing Gradient

Partial derivatives > 1


Contribution goes to
infinity

Exploding Gradient

Solving for vanishing or exploding gradients

- Identity RNN with ReLU activation
- Gradient clipping
- Skip connections

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \longrightarrow -1$$



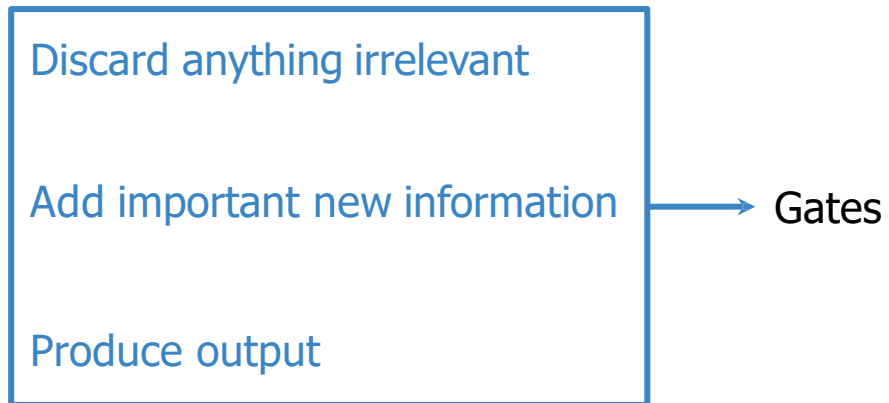
- Outline
 - Meet the Long short-term memory unit
 - LSTM architecture
 - Applications
- Memorable solution
 - Learns when to remember and when to forget
 - Basic anatomy:
 - A cell state
 - A hidden state
 - Multiple gates
 - Gates allow gradients to avoid vanishing and exploding

LSTMs: Based on previous understanding


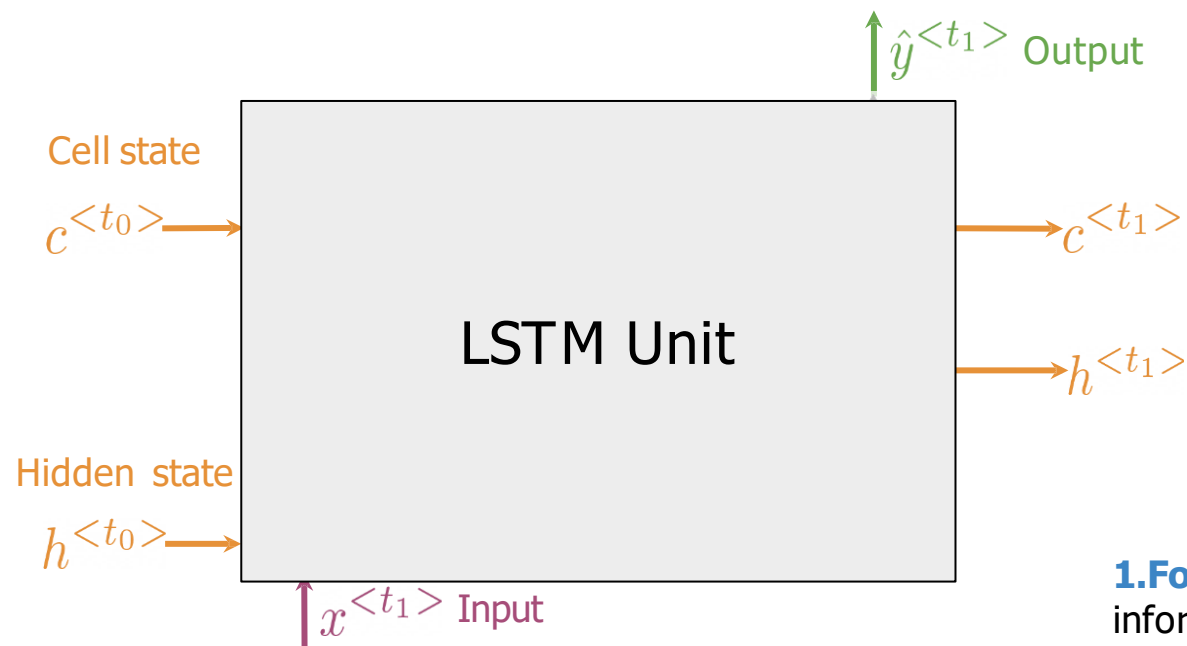
Starting point with some irrelevant information



Cell and Hidden States



Gates in LSTM



1. Forget Gate:

information that is no longer important

2. Input Gate: information to be stored

3. Output Gate:

information to use at current step

Applications of LSTMs

Next-character
prediction



Chatbots



Music
composition



Image
captioning

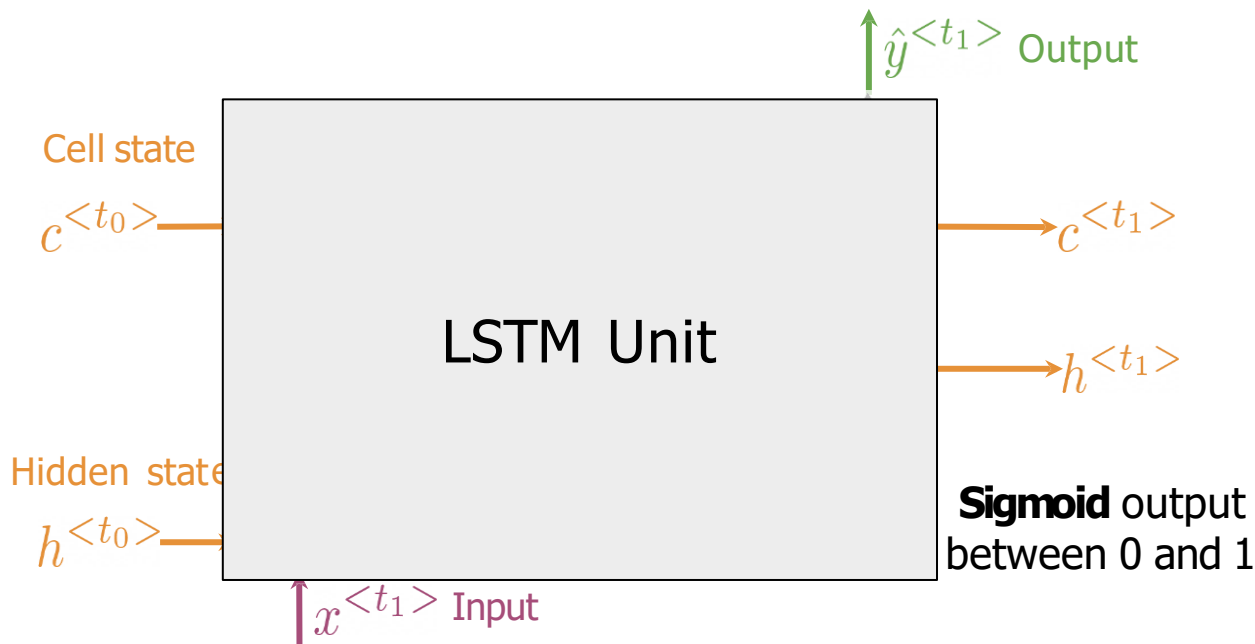


Speech
recognition



LSTM Architecture

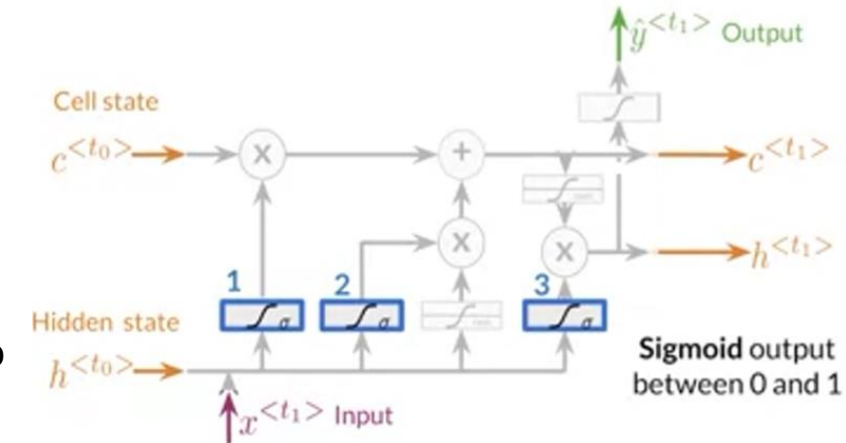
- Gates in LSTM



1. Forget Gate: information that is no longer important

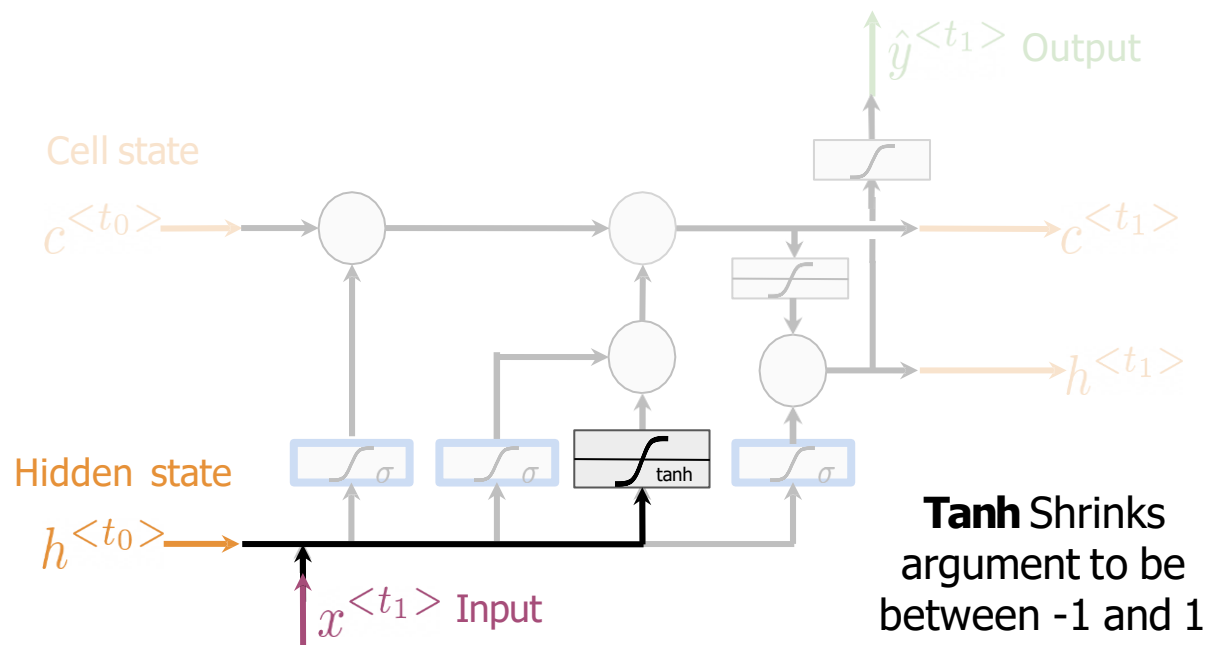
2. Input Gate: information to be stored

3. Output Gate: information to use at current step



Candidate Cell State

- Transform the information from the previous hidden states and the current inputs
- A hyperbolic tangent activation function shrinks the information to be between -1 and 1



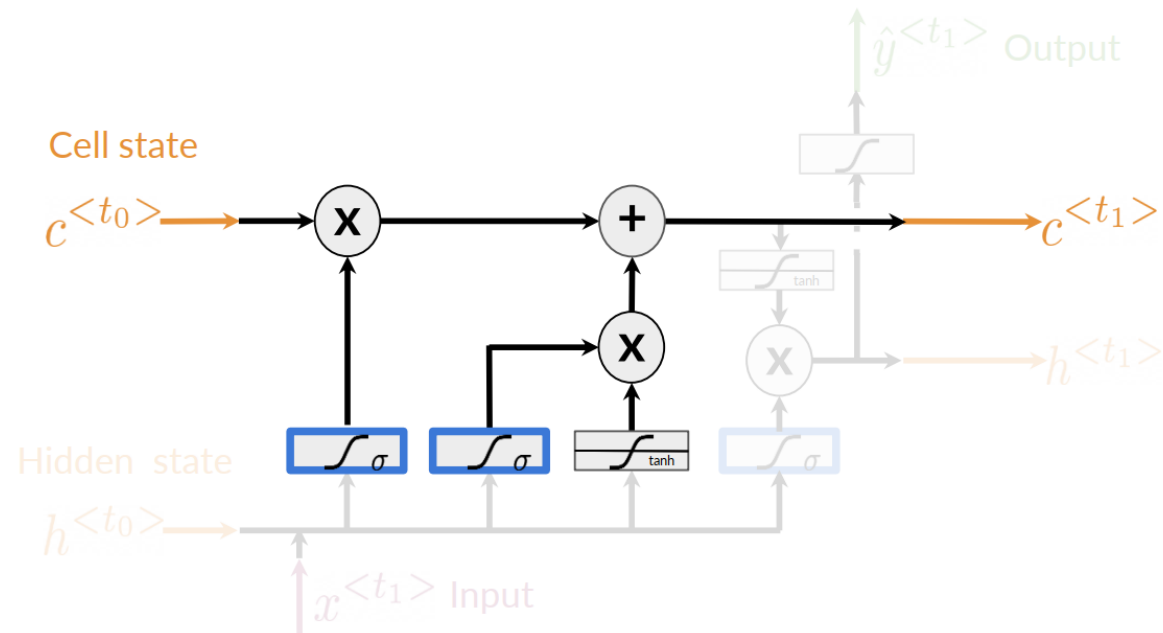
Candidate cell state

Information from the previous **hidden state** and current **input**

Tanh Shrinks argument to be between -1 and 1

Other activations could be used

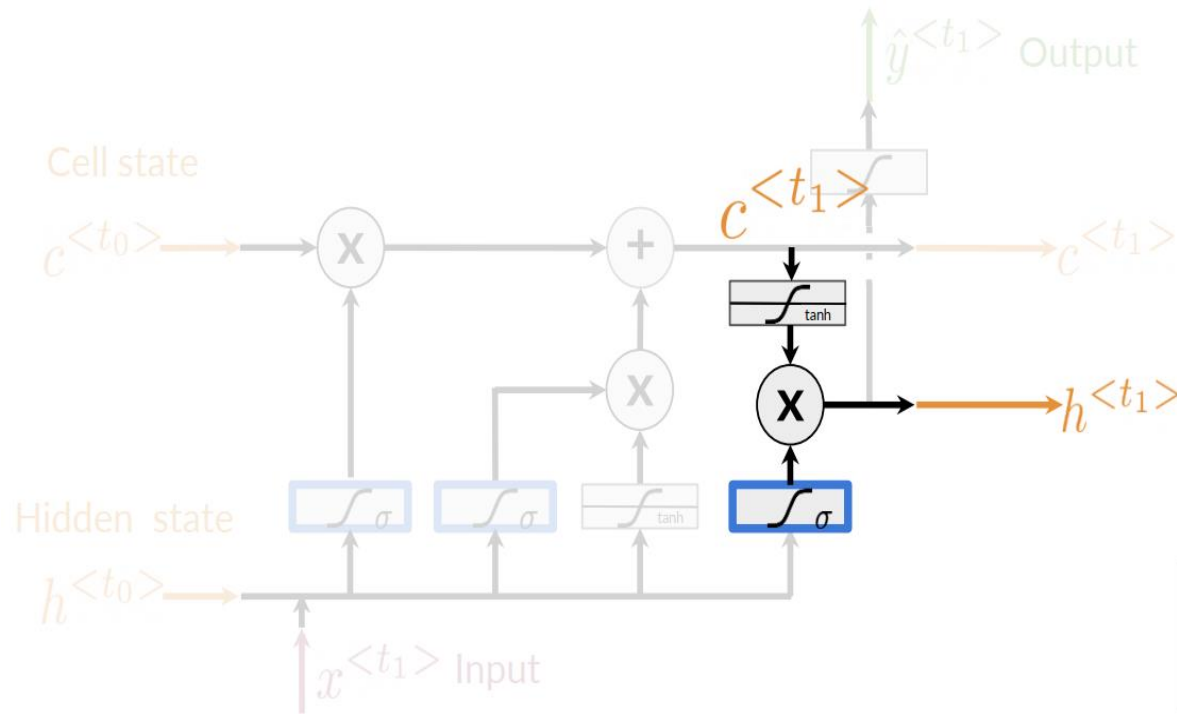
New Cell State



New Cell state

Add information from the **candidate cell state** using the **forget** and **input gates**

New Hidden State



New Hidden State

Select information from the **new cell state** using the **output gate**

The **Tanh** activation could be omitted

- LSTMs use a series of gates to decide which information to keep:
 - Forget gate decides what to keep
 - Input gate decides what to add
 - Output gate decides what the next hidden state will be

Introduction to Named Entity Recognition

- What is Named Entity Recognition?
 - Locates and extracts predefined entities from text
 - Places, organizations, names, time and dates
- Types of Entities



Thailand:
Geographical



Google:
Organization



Indian:
Geopolitical

More Types of Entities



December:
Time Indicator



Egyptian statue:
Artifact



Barack Obama:
Person

Example of a labeled sentence



- Applications of NER systems
 - Search engine efficiency
 - Recommendation engines
 - Customer service
 - Automatic trading

Training NERs: Data Processing

- Convert words and entity classes into arrays
- Token padding
- Create a data generator
- Processing data for NERs
 - Assign each class a number
 - Assign each word a number

Sharon flew to **Miami** last **Friday**.

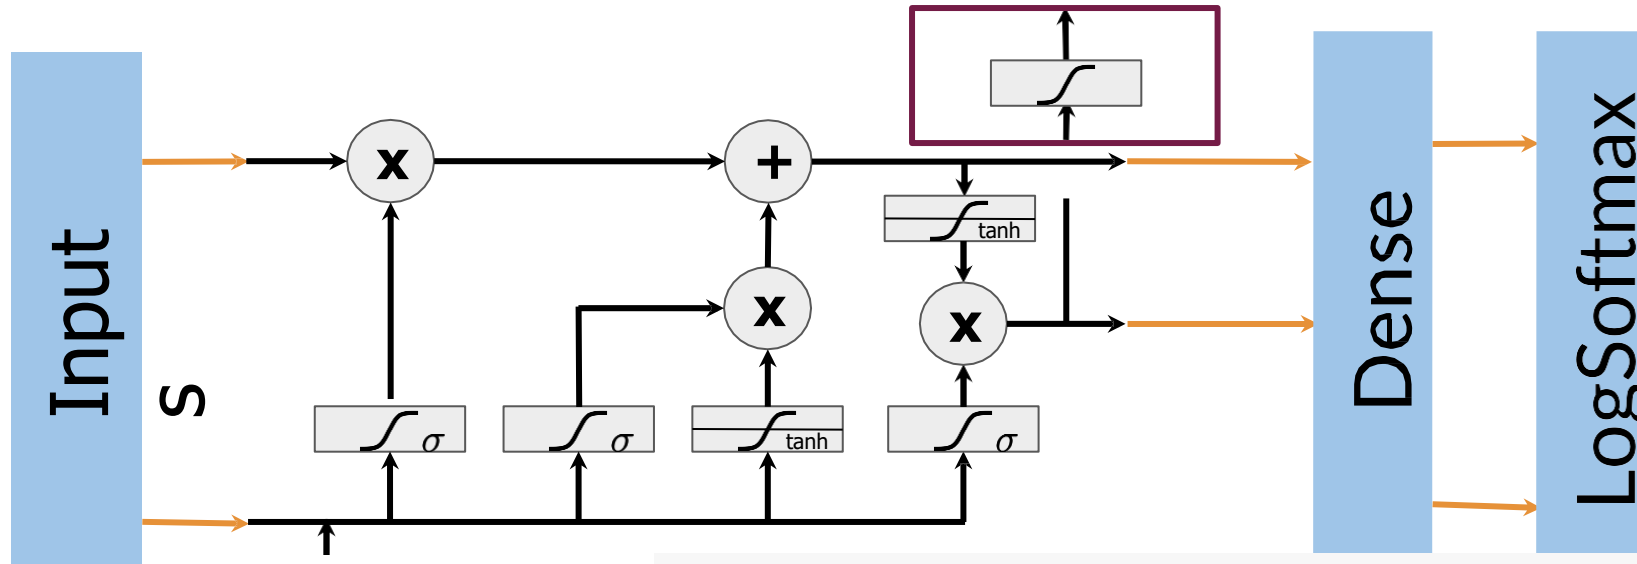
[**42821**, 853, 187, **53882**, 2894, **73**]

B-per 0 0 **B-geo** 0 **B-tim**

Token padding

- For LSTMs, all sequences need to be the same size.
- Set sequence length to a certain number
- Use the <PAD> token to fill empty spaces
- Training the NER
 - 1. Create a tensor for each input and its corresponding number
 - 2. Put them in a batch → 64, 128, 256, 512 ...
 - 3. Feed it into an LSTM unit
 - 4. Run the output through a dense layer
 - 5. Predict using a log softmax over K classes

Training the NER



- Layers in Trax

```
model = tl.Serial( tl.Embedding(), tl.LSTM(), tl.LogSoftmax())
```

- Summary

- Convert words and entities into same-length numerical arrays
- Train in batches for faster processing
- Run the output through a final layer and activation

Computing Accuracy

- Evaluating the model
 - 1. Pass test set through the model
 - 2. Get arg max across the prediction array
 - 3. Mask padded tokens
 - 4. Compare outputs against test labels

Evaluating the model in Python

```
def evaluate_model(test_sentences, test_labels, model):  
    pred = model(test_sentences)  
    outputs = np.argmax(pred, axis=2)  
    mask = ...  
    accuracy =   
    np.sum(outputs==test_labels)/float(np.sum(mask))  
  
    return accuracy
```

- If padding tokens, remember to mask them when computing accuracy